

Introduction to Time Series Analysis

Kritchavat Ploddi

Section 1

ARIMA and Exponential Smoothing Models

Packages requirement

Before proceeding, please download and install following packages:

```
library(forecast)
library(tseries)
library(ggplot2)
library(dplyr)
```

Another important concept in time series analysis is stationary.


This is because most time series forecasting models require stationary assumptions.

Stationary

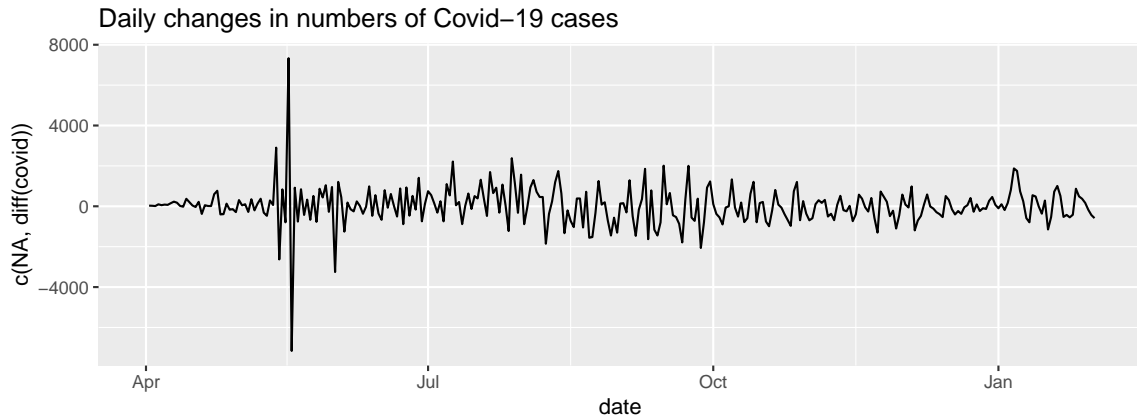
A time series Y is said to be stationary if all of its values y_t do not depend on time t .
In other words, the distribution of y_t has constant mean and variance.

Stationary ?

The numbers of Covid-19 cases



Stationary ?



How to transform non-stationary time series to be stationary

Stationary is characterized by constant mean and variance.

Transformations help to stabilize the variance.

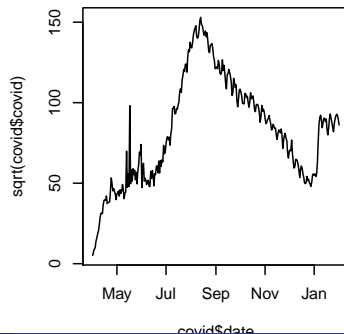

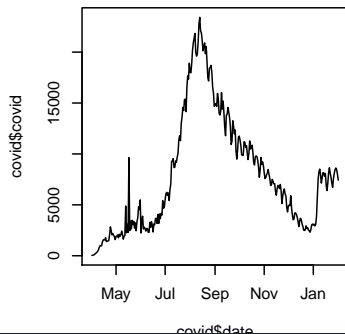
Difference help to stabilize the mean

Variance could be stabilized by taking:

- (Natural) Logarithm (\log)
- Squared root ($\sqrt{}$)

Transformations

```
par(mfrow=c(1,3))  
plot(covid$date, covid$covid, type='l')  
plot(covid$date, log(covid$covid), type='l')  
plot(covid$date, sqrt(covid$covid), type='l')
```



Transformations

We found that taking natural logarithm on Covid cases mostly stabilize the variance.

```
covid$log_covid <- log(covid$covid)
head(covid[, c('date', 'covid', 'log_covid')])
```

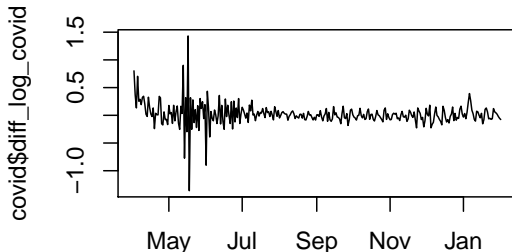

```
##           date covid log_covid
## 1 2021-04-01     26  3.258097
## 2 2021-04-02     58  4.060443
## 3 2021-04-03     84  4.430817
## 4 2021-04-04     96  4.564348
## 5 2021-04-05    194  5.267858
## 6 2021-04-06    250  5.521461
```

Differencing

Difference helps stabilize the mean.

In R, we could use the command `diff()` (plus NA offset at the first index).

```
covid$diff_log_covid <- c(NA, diff(covid$log_covid))  
  
par(mfrow=c(1,2))  
plot(covid$date, covid$covid, type='l')  
plot(covid$date, covid$diff_log_covid, type='l')
```



How could we statistically test if a series is stationary?

We could use KPSS test (Kwiatkowski–Phillips–Schmidt–Shin test) to test if a series is stationary. Using the function `kpss.test()` from `tseries` package.

```
print(kpss.test((covid$diff_log_covid))) # Log transform with differencing

##
##  KPSS Test for Level Stationarity
##
## data:  (covid$diff_log_covid)
## KPSS Level = 0.9128, Truncation lag parameter = 5, p-value = 0.01
```

How could we statistically test if a series is stationary?

This suggested that log-transformed differencing is not enough.

We have to take another differencing:

```
diff_diff_log_covid <- diff(na.omit(covid$diff_log_covid))  
print(kpss.test(diff_diff_log_covid)) # Log transform with second differencing
```

```
##  
## KPSS Test for Level Stationarity  
##  
## data: diff_diff_log_covid  
## KPSS Level = 0.068379, Truncation lag parameter = 5, p-value = 0.1
```

How could we statistically test if a series is stationary?

We could use KPSS test (Kwiatkowski–Phillips–Schmidt–Shin test) to test if a series is stationary. Using the function `kpss.test()` from `tseries` package.

```
print(kpss.test(covid$log_covid)) # Log transform
```

```
##  
## KPSS Test for Level Stationarity  
##  
## data: covid$log_covid  
## KPSS Level = 1.6925, Truncation lag parameter = 5, p-value = 0.01
```

How could we statistically test if a series is stationary?

We could use KPSS test (Kwiatkowski–Phillips–Schmidt–Shin test) to test if a series is stationary. Using the function `kpss.test()` from `tseries` package.

```
print(kpss.test(covid$covid)) # Original series
```

```
##  
## KPSS Test for Level Stationarity  
##  
## data: covid$covid  
## KPSS Level = 1.105, Truncation lag parameter = 5, p-value = 0.01
```


Autoregressive Integrated Moving Average (ARIMA) Models

An ARIMA model is rarely interpretable in terms of visible data structures like trend and seasonality. But it can capture a huge range of time series patterns.

Autoregressive Integrated Moving Average (ARIMA) Models

An ARIMA model is rarely interpretable in terms of visible data structures like trend and seasonality. But it can capture a huge range of time series patterns.

It is composed of an Autoregressive (AR) and Moving Average (MA) models

Autoregressive models (AR)

$$y_t = c + \phi_1 y_{t-1} + \phi_2 y_{t-2} + \cdots + \phi_p y_{t-p} + \varepsilon_t,$$

This is called AR(p) model.

AR models predict future values based on previous values (at p lags).

Moving Average Models (MA)

$$y_t = c + \varepsilon_t + \theta_1 \varepsilon_{t-1} + \theta_2 \varepsilon_{t-2} + \cdots + \theta_q \varepsilon_{t-q},$$

This is called MA(q) model.

MA models predict future values based on previous errors (at q lags).

ARMA and ARIMA models

ARMA(p, q) is a linear combination of AR(p) and MA(q). ARMA predict future values based on both previous values and errors.

ARIMA(p, n, q) is similar to ARMA(p, q), except only n difference has been taken. This is done for taking an account for non-stationary time series.

Forecasting Dengue with ARIMA models

```
dhfrain <- read.csv('data/dhfrain2.csv')
dhfrain$date <- as.Date(dhfrain$date)
dhfrain[, c("date", "cases")]
```

##		date	cases
## 1		2009-01-01	2614
## 2		2009-02-01	2057
## 3		2009-03-01	2324
## 4		2009-04-01	2947
## 5		2009-05-01	6234
## 6		2009-06-01	8569
## 7		2009-07-01	7184
## 8		2009-08-01	7302
## 9		2009-09-01	5016
## 10		2009-10-01	4640
## 11		2009-11-01	4723

Forecasting Dengue with ARIMA models


Train-test splitting

```
train <- dhfrain[dhfrain$year < 2018, ]  
test  <- dhfrain[dhfrain$year == 2018, ]
```

Train-test splitting




Cross-validation




Forecasting Dengue with ARIMA models

```
ggplot(train) + geom_line(aes(x=date, y=cases))
```



Forecasting Dengue with ARIMA models

```
train$logcases <- log(train$cases)
logcases_diff <- diff(train$logcases)
plot(logcases_diff, type='l')
```



Forecasting Dengue with ARIMA models

KPSS test for stationary

```
library(tseries)
kpss.test(logcases_diff)
```


```
##
##  KPSS Test for Level Stationarity
##
## data:  logcases_diff
## KPSS Level = 0.038409, Truncation lag parameter = 4, p-value = 0.1
```

Forecasting Dengue with ARIMA models


ACF and PACF plot

```
par(mfrow=c(1,2))  
acf(logcases_diff)  
pacf(logcases_diff)
```

Series logcases_diff



Series logcases_diff




Forecasting Dengue with ARIMA models


ACF and PACF plot (seasonal)

```
par(mfrow=c(1,2))  
acf(logcases_diff, lag=60)  
pacf(logcases_diff, lag=60)
```

Series logcases_diff



Series logcases_diff



Forecasting Dengue with ARIMA models

So the ARIMA model would be: (S)ARIMA (1,1,0) (2,1,0), 12

```
manual <- arima(train$logcases,  
                order=c(1,1,0),  
                seasonal=list(order=c(3,1,0), period=12))
```

Auto ARIMA

```
library(forecast)
auto <- auto.arima(train$logcases)
```


Models evaluation

```
test$forecast_manual <- exp(forecast(manual, h=12)$mean)
test$forecast_auto <- exp(forecast(auto, h=12)$mean)
```

Models evaluation: metrics

Error $Error = Y_{actual} - Y_{forecast}$

Absolute Error $AbsError = Abs(Error)$

Absolute Percentage Error $AbsPctError = Abs(Error) / Y_{actual} \times 100$

Squared Error $SqError = Error^2$

Mean Squared Error $MSE = mean(SqError)$


Root mean squared error $RMSE = \sqrt{MSE}$

Mean Absolute Error $MAE = mean(AbsError)$

Mean Absolute Percentage Error $MAPE = mean(AbsPctError)$

Models evaluation

```
ggplot() +  
  geom_line(data = train, aes(x=date, y=cases)) +  
  geom_line(data = test, aes(x=date, y=cases)) +  
  geom_line(data = test, aes(x=date, y=forecast_auto), color='red') +  
  geom_line(data = test, aes(x=date, y=forecast_manual), color='blue')
```



Models evaluation

Root mean squared error:

```
error_manual <- with(test, forecast_manual - cases)
error_auto <- with(test, forecast_auto - cases)
```

```
# SE
```

```
sqerror_manual <- error_manual^2
sqerror_auto <- error_auto^2
```

```
# MSE
```

```
mse_manual <- mean(sqerror_manual)
mse_auto <- mean(sqerror_auto)
```

```
# RMSE
```

```
rmse_manual <- sqrt(mse_manual)
rmse_auto <- sqrt(mse_auto)
```

```
c(manual=rmse_manual, auto=rmse_auto)
```

```
##      manual      auto
## 4298.099 3096.161
```

Models evaluation

Mean absolute error: (MAE)

```
error_manual <- with(test, forecast_manual - cases)
error_auto <- with(test, forecast_auto - cases)
```

```
# ABS Error
```

```
abserror_manual <- abs(error_manual)
abserror_auto <- abs(error_auto)
```

```
# MAE
```

```
mae_manual <- mean(abserror_manual)
mae_auto <- mean(abserror_auto)
```

```
c(manual=mae_manual, auto=mae_auto)
```

```
## manual auto
## 3234.929 2094.105
```

Model evaluation

Mean absolute percentage error: (MAE)

```
error_manual <- with(test, forecast_manual - cases)
error_auto <- with(test, forecast_auto - cases)

# ABS Error

abserror_manual <- abs(error_manual)
abserror_auto <- abs(error_auto)

# Abs Percentage Error
abs_pct_error_manual <- abserror_manual/test$cases * 100
abs_pct_error_auto <- abserror_auto/test$cases * 100

# MAPE
mape_manual <- mean(abs_pct_error_manual)
mape_auto <- mean(abs_pct_error_auto)

c(manual=mape_manual, auto=mape_auto)

##      manual      auto
## 38.90675 23.20156
```

Section 2

Exponential smoothing methods

Naïve average

Using the naïve method, all forecasts for the future are equal to the last observed value of the series,


$$\hat{y}_{T+h|T} = y_T$$

##		date	rti	Naïve
## 1		2011-01-01	2068	NA
## 2		2011-02-01	1859	2068
## 3		2011-03-01	1987	1859
## 4		2011-04-01	1907	1987
## 5		2011-05-01	1763	1907
## 6		2011-06-01	1678	1763
## 7		2011-07-01	1709	1678
## 8		2011-08-01	1706	1709
## 9		2011-09-01	1536	1706
## 10		2011-10-01	1779	1536
## 11		2011-11-01	1925	1779
## 12		2011-12-01	2079	1925
## 13		2012-01-01	1968	2079
## 14		2012-02-01	1840	1968
## 15		2012-03-01	2109	1840
## 16		2012-04-01	1963	2109
## 17		2012-05-01	1697	1963
## 18		2012-06-01	1680	1697
## 19		2012-07-01	1724	1680
## 20		2012-08-01	1597	1724

Average method

Using the average method, all future forecasts are equal to a simple average of the observed data,

$$\hat{y}_{T+h|T} = \frac{1}{T} \sum_{t=1}^T y_t,$$



Simple Exponential Smoothing

$$\hat{y}_{T+1|T} = \alpha y_T + \alpha(1 - \alpha)y_{T-1} + \alpha(1 - \alpha)^2 y_{T-2} + \cdots$$

Simple Exponential Smoothing


```
rti <- read.csv("data/rti.csv")
rti$date <- as.Date(rti$date)
ts_rti <- ts(rti$rti, start=c(2011,1), freq=12)
ts_rti
```

##		Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
##	2011	2068	1859	1987	1907	1763	1678	1709	1706	1536	1779	1925	2079
##	2012	1968	1840	2109	1963	1697	1680	1724	1597	1496	1723	1853	1953
##	2013	1936	1767	1983	1856	1717	1671	1657	1540	1502	1675	1909	2008
##	2014	2091	1897	2080	1764	1682	1565	1459	1493	1410	1644	1730	1972
##	2015	1786	1596	1726	1799	1574	1499	1477	1456	1473	1672	1860	2041
##	2016	2023	1956	1999	1973	1645	1661	1686	1554	1598	1708	1860	2082
##	2017	2018	1902	2036	1827	1688	1723	1509	1568	1671	1742	1827	2096
##	2018	1841	1772	1904	1763	1593	1470	1504	1469	1437	1550	1788	1839
##	2019	1908	1758	1853	1781	1570	1584	1548	1443	1393	1570	1662	1834
##	2020	1679	1628	1614	916	1296	1281	1424	1452	1388	1602	1707	1844

Simple Exponential Smoothing

```
forecast_rti <- ses(ts_rti, h=15)  
autoplot(forecast_rti) +  
  autolayer(fitted(forecast_rti), series="Fitted")
```


Forecasts from Simple exponential smoothing



Additive and multiplicative trend

Additive trend


```
data.frame(t=1:100, e=rnorm(100,0,5)) %>%  
  mutate(y=20+1.6*t+e) %>%  
  ggplot() + geom_line(aes(x=t, y=y)) + ylim(0,300)
```



Additive and multiplicative trend

Multiplicative trend

```
data.frame(t=1:100, e=rnorm(100,0,10)) %>%  
  mutate(y=exp(0.06*t)+e) %>%  
  ggplot() + geom_line(aes(x=t, y=y))
```





Additive and multiplicative seasonal



Seasonality

Additive

Multiplicative



Trend



Classification of exponential smoothing method

Trend Component	Seasonal Component		
	N	A	M
	(None)	(Additive)	(Multiplicative)
N (None)	(N,N)	(N,A)	(N,M)
A (Additive)	(A,N)	(A,A)	(A,M)
A_d (Additive damped)	(A_d ,N)	(A_d ,A)	(A_d ,M)

Some of these methods we have already seen using other names:

Short hand	Method
(N,N)	Simple exponential smoothing
(A,N)	Holt's linear method
(A_d ,N)	Additive damped trend method
(A,A)	Additive Holt-Winters' method
(A,M)	Multiplicative Holt-Winters' method
(A_d ,M)	Holt-Winters' damped method

Fitting exponential smoothing method in R

In R, using `ets()` function from `forecast` package to fit an exponential smoothing


We can specify the type of `ets` using the model parameters.

- “NNN” Simple Average
- “NAN” Additive Trend without seasonal
- “NMN” Multiplicative Trend without seasonal
- “NAA” Additive Trend and seasonal (Additive Holt-Winter’s)
- “NMM” Multiplicative Trend and seasonal (Multiplicative Holt-Winter’s)

We can also use “Z”, e.g. “ZZZ” to let R to find the best fitted ETS models.

Fitting exponential smoothing method in R

For example, if we wish to fit Additive Holt-Winter methods (Additive trend and seasonality) and Multiplicative Holt-Winter methods (Multiplicative trend and seasonality)




Fitting exponential smoothing method in R

Fitting Additive Holt-Winter methods (Additive trend and seasonality), "ZAA"

Fitting Multiplicative Holt-Winter methods (Multiplicative trend and seasonality), "ZMM"

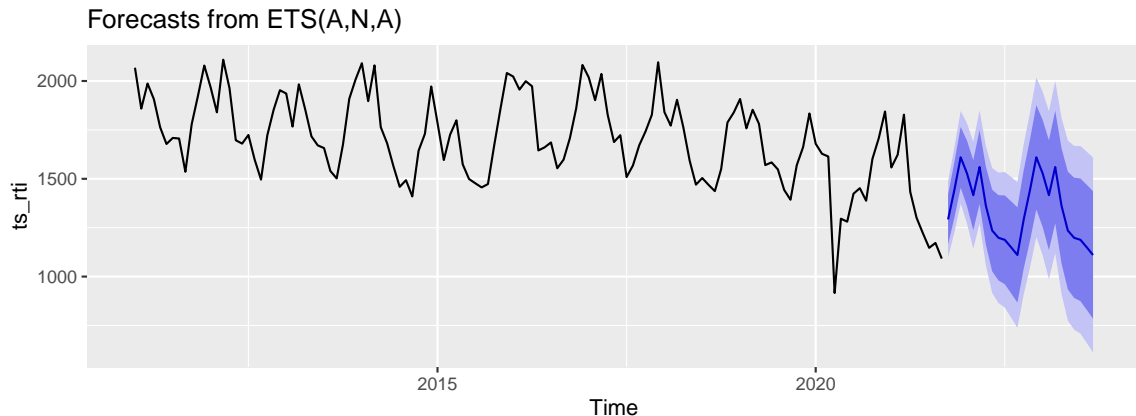
```
zaa_rti <- ets(ts_rti, model="ZAA") # Additive  
zmm_rti <- ets(ts_rti, model="ZMM") # Multiplicative  
par(mfrow=c(1,2))  
autoplot(forecast(zaa_rti), h=16)
```



Fitting exponential smoothing method in R

We can also use “ZZZ” to let R automatically fit and Exponential Smoothing Models


```
zzz_rti <- ets(ts_rti, model="ZZZ")  
autoplot(forecast(zzz_rti), h=16)
```



Fitting an ETS with Dengue data

We would be manually fitting an ETS with only multiplicative seasonal with training data of Dengue (2009-2017).

We would also use an automated methods.




Fitting an ETS with Dengue data

Using “ZNM” for multiplicative seasonal ETS

Using “ZZZ” for automated ETS

```
train_ts <- ts(train$cases, start=c(2009,1), freq=12)
znm_train <- ets(train_ts, model="ZNM")
zzz_train <- ets(train_ts, model="ZZZ") # Auto
par(mfrow=c(1,2))
autoplot(forecast(znm_train), h=12)
```

Forecasts from ETS(M,N,M)



Mean Absolute Percentage Errors

```
forecast_test_znm <- forecast(znm_train, h=12)$mean
forecast_test_zzz <- forecast(zzz_train, h=12)$mean

abs_error_znm <- abs(forecast_test_znm - test$cases)
abs_error_zzz <- abs(forecast_test_zzz - test$cases)

abs_pct_error_znm <- abs_error_znm/test$cases * 100
abs_pct_error_zzz <- abs_error_zzz/test$cases * 100

mape_znm <- mean(abs_pct_error_znm)
mape_zzz <- mean(abs_pct_error_zzz)

c(ZNM=mape_znm, AUTO=mape_zzz)
```

```
##      ZNM      AUTO
## 34.62963 34.62963
```

Models evaluation

Comparing MAPE with previous ARIMA models

```
c(AUTO_ETS=mape_zzz, manual_ARIMA=mape_manual, AUTO_ARIMA=mape_auto)
```

##	AUTO_ETS	manual_ARIMA	AUTO_ARIMA
##	34.62963	38.90675	23.20156